

Fine-tuning leaves a recoverable family history in model weights.


A. Can weights recover the tree?



✓ every true branch recovered in 85–100% of runs

(3% for a random tree); one attention key matrix,
0.1–0.2% of the weights, suffices

B. Do weights track behavior?



C. When does recovery work?

